

# Predictive Modeling

## Series 7

### Exercise 7.1

In this problem, you will use support vector approaches in order to predict whether a given car gets high or low gas mileage based on the **Auto** data set.

- Create a binary variable that takes on a 1 for cars with gas mileage above the median, and a 0 for cars with gas mileage below the median.
- Fit a support vector classifier to the data with various values of `cost`, in order to predict whether a car gets high or low gas mileage. Report the cross-validation errors associated with different values of this parameter. Comment on your results.
- Now repeat (b), this time using SVMs with radial and polynomial basis kernels, with different values of `gamma` and `degree` and `cost`. Comment on your results
- Make some plots to back up your assertions in (b) and (c).

*Hint:* We used the `plot()` function for `svm` objects only in cases with  $p = 2$ . When  $p > 2$ , you can use the `plot()` function to create plots displaying pairs of variables at a time. Essentially, instead of typing

```
plot(svmfit, dat)
```

where `svmfit` contains your fitted model and `dat` is a data frame containing your data, you can type

```
plot(svmfit, dat, x1 ~ x4)
```

in order to plot just the first and fourth variables. However, you must replace `x1` and `x4` with the correct variable names. To find out more, type `?plot.svm`.

### Exercise 7.2

This problem involves the **OJ** data set, which is part of the ISLR package.

- Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

- b) Fit a support vector classifier to the training data using `cost=0.01`, with `Purchase` as the response and the other variables as predictors. Use the `summary()` function to produce summary statistics, and describe the results obtained.
- c) What are the training and test error rates? Also give the corresponding confusion matrices.
- d) Use the `tune()` function to select an optimal cost. Consider 10 values in the range 0.01 to 10.
- e) Compute the training and test error rates using this new value for `cost`.
- f) Repeat parts (b) through (e) using a support vector machine with a radial kernel. Use the default value for `gamma`.
- g) Repeat parts (b) through (e) using a support vector machine with a polynomial kernel. Set `degree=2`.
- h) Overall, which approach seems to give the best results on this data?

## Result Checker

# Predictive Modeling

## Solutions to Series 7

### Solution 7.1

a) `require` (ISLR)

```
## Loading required package: ISLR

Auto$mpg <- ifelse(Auto$mpg > median(Auto$mpg), 1, 0)
table(Auto$mpg)

##
##      0      1
## 196 196
```

b) The cross-validation errors can be produced with ease by the `tune` function from the `e1071` library. This is done by the code below.

```
require(e1071)

## Loading required package: e1071


# tuning grid for the cost parameter
costs <- data.frame(cost = seq(0.05, 100, length.out = 6))

# 10-fold cross validation
svm.tune.linear = tune(svm, mpg ~ ., data = Auto, ranges = costs,
  kernel = "linear")
svm.tune.linear

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.05
##
## - best performance: 0.103031
```

If we plot the values of the missclassification error, the minimum missclassification error occurs at cost 0.05.

```
plot(svm.tune.linear$performance[, c(1, 2)], type = "l")
```




c) Radial Kernel:

```
params <- data.frame(cost = seq(0.001, 240, length.out = 6),  
  gamma = c(0.5, 1, 2, 3, 4, 5))  
svm.tune.radial <- tune(svm, mpg ~ ., data = Auto, ranges = params,  
  kernel = "radial")  
svm.tune.radial  
  
##  
## Parameter tuning of 'svm':  
##
```

```
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
## 48.0008    0.5
##
## - best performance: 0.07764927
```

```
plot(svm.tune.radial$performance[, c(2, 3)], type = "l")
```




Polynomial Kernel:

```
params <- data.frame(cost = seq(0.001, 240, length.out = 6),
  degree = c(0.5, 1, 2, 3, 4, 5))
svm.tune.poly <- tune(svm, mpg ~ ., data = Auto, ranges = params,
  kernel = "polynomial")
svm.tune.poly

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##      cost degree
## 144.0004      1
##
## - best performance: 0.09806219
```

```
plot(svm.tune.poly$performance[, c(2, 3)], type = "l")
```



d) `plot(x = svm.tune.radial$best.model, data = Auto, formula = weight ~ horsepower)`

```
plot(x = svm.tune.radial$best.model, data = Auto, formula = cylinders ~ year)
```

## Solution 7.2



```

a) set.seed(42)
train = sample(1:1070, 800)
test = (1:1070)[-train]

tb = c()
res = c()

b) require(ISLR)
require(e1071)

svm.fit = svm(Purchase ~ ., data = OJ, subset = train, cost = 0.01,
              kernel = "linear")
summary(svm.fit)

##
## Call:
## svm(formula = Purchase ~ ., data = OJ, cost = 0.01,
##      kernel = "linear", subset = train)
##
##
## Parameters:
##      SVM-Type:  C-classification
##      SVM-Kernel: linear
##              cost: 0.01
##
## Number of Support Vectors: 432
##
##      ( 215 217 )
##
##
## Number of Classes: 2
##
## Levels:
##      CH MM

```

From the output of **R**'s summary function we can see that 432 observations are used as support vector. Moreover, the support vectors are almost equally split among the classes.

```

c) # train
svm.pred = predict(svm.fit, OJ[train, ])
kable(table(OJ[train, "Purchase"], svm.pred))

```

	CH	MM
CH	432	60
MM	77	231

```
mean(OJ$Purchase[train] != svm.pred)
```

```
## [1] 0.17125
```

```
res = cbind(res, train = mean(OJ$Purchase[train] != svm.pred))
```

```
# test
```

```
svm.pred = predict(svm.fit, OJ[test, ])
```

```
kable(table(OJ[test, "Purchase"], svm.pred))
```

	CH	MM
CH	142	19
MM	25	84

```
mean(OJ$Purchase[test] != svm.pred)
```

```
## [1] 0.162963
```

```
res = cbind(res, test = mean(OJ$Purchase[test] != svm.pred))
```

- d) `svm.tune = tune(svm, Purchase ~ ., data = OJ[train, ], ranges = data.frame(10, 25)), kernel = "linear")`  
`summary(svm.tune)`

```
##
```

```
## Error estimation of 'svm' using 10-fold cross validation: 0.1825
```

```
res = cbind(res, CV = svm.tune$best.performance)
```

- e) `svm.pred = predict(svm.tune$best.model, OJ[train, ])`  
`kable(table(OJ[train, "Purchase"], svm.pred))`

	CH	MM
CH	432	60
MM	77	231

```
mean(OJ$Purchase[train] != svm.pred)
```

```
## [1] 0.17125
```

```
res = cbind(res, train.tuned = mean(OJ$Purchase[train] !=  
svm.pred))
```

```
svm.pred = predict(svm.tune$best.model, OJ[test, ])
```

```
kable(table(OJ[test, "Purchase"], svm.pred))
```

	CH	MM
CH	142	19
MM	25	84

```
mean(OJ$Purchase[test] != svm.pred)
```

```
## [1] 0.162963
```

```
res = cbind(res, test.tuned = mean(OJ$Purchase[test] !=
  svm.pred))
```

```
tb = rbind(tb, res)
```

```
res = c()
```

```
f) # b
svm.fit = svm(Purchase ~ ., data = OJ, subset = train, cost = 0.01,
  kernel = "radial")
summary(svm.fit)

##
## Call:
## svm(formula = Purchase ~ ., data = OJ, cost = 0.01,
##     kernel = "radial", subset = train)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##     cost:  0.01
##
## Number of Support Vectors:  621
##
##   ( 308 313 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM

# train
svm.pred = predict(svm.fit, OJ[train, ])
kable(table(OJ[train, "Purchase"], svm.pred))
```

	CH	MM
CH	492	0
MM	308	0

```
mean(OJ$Purchase[train] != svm.pred)
```

```
## [1] 0.385
```

```
res = cbind(res, train = mean(OJ$Purchase[train] != svm.pred))
```

```
# test
```

```
svm.pred = predict(svm.fit, OJ[test, ])
```

```
kable(table(OJ[test, "Purchase"], svm.pred))
```

	CH	MM
CH	161	0
MM	109	0

```
mean(OJ$Purchase[test] != svm.pred)
```

```
## [1] 0.4037037
```

```
res = cbind(res, train = mean(OJ$Purchase[test] != svm.pred))
```

```
svm.tune = tune(svm, Purchase ~ ., data = OJ[train, ], ranges = data.frame(
  10, 25))
```

```
summary(svm.tune)
```

```
##
```

```
## Error estimation of 'svm' using 10-fold cross validation: 0.385
```

```
res = cbind(res, CV = svm.tune$best.performance)
```

```
# train
```

```
svm.pred = predict(svm.tune$best.model, OJ[train, ])
```

```
kable(table(OJ[train, "Purchase"], svm.pred))
```

	CH	MM
CH	492	0
MM	308	0

```
mean(OJ$Purchase[train] != svm.pred)
```

```
## [1] 0.385
```

```
res = cbind(res, train.tuned = mean(OJ$Purchase[train] !=
  svm.pred))
```

```
# test
```

```
svm.pred = predict(svm.tune$best.model, OJ[test, ])
kable(table(OJ[test, "Purchase"], svm.pred))
```

	CH	MM
CH	161	0
MM	109	0

```
mean(OJ$Purchase[test] != svm.pred)
```

```
## [1] 0.4037037
```

```
res = cbind(res, test.tuned = mean(OJ$Purchase[test] !=
  svm.pred))
```

```
tb = rbind(tb, res)
```

```
res = c()
```

g) # b

```
svm.fit = svm(Purchase ~ ., data = OJ, subset = train, cost = 0.01,
  degree = 2, kernel = "polynomial")
summary(svm.fit)
```

```
##
```

```
## Call:
```

```
## svm(formula = Purchase ~ ., data = OJ, cost = 0.01,
##      degree = 2, kernel = "polynomial", subset = train)
```

```
##
```

```
##
```

```
## Parameters:
```

```
##      SVM-Type:  C-classification
```

```
##      SVM-Kernel: polynomial
```

```
##      cost: 0.01
```

```
##      degree: 2
```

```
##      coef.0: 0
```

```
##
```

```
## Number of Support Vectors: 621
```

```
##
```

```
##      ( 308 313 )
```

```
##
```

```
##
## Number of Classes: 2
##
## Levels:
## CH MM

# train
svm.pred = predict(svm.fit, OJ[train, ])
kable(table(OJ[train, "Purchase"], svm.pred))
```

	CH	MM
CH	492	0
MM	308	0

```
mean(OJ$Purchase[train] != svm.pred)

## [1] 0.385

res = cbind(res, train = mean(OJ$Purchase[train] != svm.pred))

# test
svm.pred = predict(svm.fit, OJ[test, ])
kable(table(OJ[test, "Purchase"], svm.pred))
```

	CH	MM
CH	161	0
MM	109	0

```
mean(OJ$Purchase[test] != svm.pred)

## [1] 0.4037037

res = cbind(res, test = mean(OJ$Purchase[test] != svm.pred))

svm.tune = tune(svm, Purchase ~ ., data = OJ[train, ], ranges = data.frame(
  10, 25)), kernel = "polynomial", degree = 2)
summary(svm.tune)

##
## Error estimation of 'svm' using 10-fold cross validation: 0.385

res = cbind(res, CV = svm.tune$best.performance)

# train
svm.pred = predict(svm.tune$best.model, OJ[train, ])
kable(table(OJ[train, "Purchase"], svm.pred))
```

	CH	MM
CH	492	0
MM	308	0

```
mean(OJ$Purchase[train] != svm.pred)
```

```
## [1] 0.385
```

```
res = cbind(res, train.tuned = mean(OJ$Purchase[train] !=
  svm.pred))
```

```
# test
```

```
svm.pred = predict(svm.tune$best.model, OJ[test, ])
```

```
kable(table(OJ[test, "Purchase"], svm.pred))
```

	CH	MM
CH	161	0
MM	109	0

```
mean(OJ$Purchase[test] != svm.pred)
```

```
## [1] 0.4037037
```

```
res = cbind(res, test.tuned = mean(OJ$Purchase[test] !=
  svm.pred))
```

```
tb = rbind(tb, res)
```

- h) The linear SVM performs best on this data. However, further investigating the optimal parameters for **gamma**, **degree** and **cost** could improve the behaviour of different classifiers.

```
rownames(tb) = c("LINEAR", "POLYNOMIAL", "RADIAL")
```

```
kable(tb)
```

	train	test	CV	train.tuned	test.tuned
LINEAR	0.17125	0.1629630	0.1825	0.17125	0.1629630
POLYNOMIAL	0.38500	0.4037037	0.3850	0.38500	0.4037037
RADIAL	0.38500	0.4037037	0.3850	0.38500	0.4037037